

```
In [ ]: import pandas as pd
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.metrics import (accuracy_score, precision_score, recall_score, f1_score,
                             confusion_matrix, roc_auc_score, classification_report, roc_curve, auc)
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import (RandomForestClassifier, GradientBoostingClassifier,
                              AdaBoostClassifier, BaggingClassifier, VotingClassifier)

from sklearn.svm import SVC
from sklearn.neighbors import KNeighborsClassifier
from sklearn.naive_bayes import GaussianNB
import xgboost as xgb
import matplotlib.pyplot as plt
import seaborn as sns
from imblearn.over_sampling import SMOTE
```

```
In [2]: # Load the dataset
file_path = 'Strokesdataset.csv'
stroke_data = pd.read_csv(file_path)
```

```
In [3]: # Drop the 'id' column
stroke_data.drop(columns=['id'], inplace=True)
```

```
In [4]: # Checking for null values
stroke_data.isnull().sum()
```

```
Out[4]: gender                0
age                0
hypertension       0
heart_disease      0
ever_married       0
work_type          0
Residence_type     0
avg_glucose_level  0
bmi                1462
smoking_status     13292
stroke             0
dtype: int64
```

```
In [5]: # Checking for duplicates
stroke_data.duplicated().sum()
```

```
Out[5]: 0
```

```
In [ ]: # Handling missing values in the BMI column by filling with the median BMI value
stroke_data['bmi'].fillna(stroke_data['bmi'].median(), inplace=True)
stroke_data['smoking_status'].fillna(stroke_data['smoking_status'].mode()[0],
                                     inplace=True)
```

```
In [ ]: # Encoding categorical variables
label_encoders = {}
for column in ['gender', 'ever_married', 'work_type',
               'Residence_type', 'smoking_status']:
    le = LabelEncoder()
    stroke_data[column] = le.fit_transform(stroke_data[column])
    label_encoders[column] = le
```

```
In [8]: stroke_data['smoking_status']
```

```
Out[8]: 0      1
        1      1
```

```

2         1
3         0
4         1
..
43395     1
43396     0
43397     0
43398     1
43399     1
Name: smoking_status, Length: 43400, dtype: int32

```

```

In [ ]: # Scaling numerical features
scaler = StandardScaler()
stroke_data[['age', 'avg_glucose_level', 'bmi']] = scaler.fit_transform(stroke_data[['ag
                                                    'avg_glucose_level', 'bm

```

```

In [10]: # Splitting the data into features and target variable
X = stroke_data.drop(columns=['stroke'])
y = stroke_data['stroke']

```

```

In [11]: # Count the occurrences of each class before SMOTE
class_counts_before = y.value_counts()

```

```

In [12]: # Apply SMOTE to balance the dataset
smote = SMOTE(random_state=42)
X_resampled, y_resampled = smote.fit_resample(X, y)

```

```

In [13]: # Count the occurrences of each class after SMOTE
class_counts_after = y_resampled.value_counts()

```

```

In [14]: # Create a DataFrame to hold the class counts
class_counts_df = pd.DataFrame({
    'Before SMOTE': class_counts_before,
    'After SMOTE': class_counts_after
}).reset_index()

```

```

In [15]: class_counts_df.columns = ['Class', 'Before SMOTE', 'After SMOTE']


```

```

In [ ]: # Plot for class counts before SMOTE
plt.figure(figsize=(8, 6))
sns.barplot(x='Class', y='Before SMOTE', data=class_counts_df,
            hue='Class', palette=['#f4c145', '#5752d1'], dodge=False)
plt.title('Class Distribution Before SMOTE')
plt.ylabel('Count')
plt.xlabel('Class')
plt.legend(title='Class')
plt.tight_layout()
plt.show()


```

Class Distribution Before SMOTE



```
In [ ]: # Plot for class counts after SMOTE
plt.figure(figsize=(8, 6))
sns.barplot(x='Class', y='After SMOTE', data=class_counts_df,
            hue='Class', palette=["#4cb140", "#009596"], dodge=False)
plt.title('Class Distribution After SMOTE')
plt.ylabel('Count')
plt.xlabel('Class')
plt.legend(title='Class')
plt.tight_layout()
plt.show()
```

Class Distribution After SMOTE



```
In [ ]: # Splitting the resampled data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X_resampled, y_resampled,
                                                    test_size=0.2, random_state=42)
```

```
In [ ]: # Dictionary to store the models and their results
models = {
    'Logistic Regression': LogisticRegression(max_iter=1000),
    'Decision Tree': DecisionTreeClassifier(),
    'Random Forest': RandomForestClassifier(),
    'Gradient Boosting': GradientBoostingClassifier(),
    'AdaBoost': AdaBoostClassifier(),
    'XGBoost': xgb.XGBClassifier(eval_metric='logloss'),
    'Support Vector Machine': SVC(probability=True),
    'K-Nearest Neighbors': KNeighborsClassifier(),
    'Naive Bayes': GaussianNB(),
    'Bagging': BaggingClassifier(estimator=DecisionTreeClassifier(),
                                n_estimators=50, random_state=42),
    'Voting': VotingClassifier(estimators=[
        ('lr', LogisticRegression(max_iter=1000)),
        ('rf', RandomForestClassifier()),
        ('svc', SVC(probability=True))
    ], voting='soft')
}
```

```
In [20]: results = []
roc_curves = {}
```

```
In [21]: # Training and evaluating the models
for model_name, model in models.items():
    print(f"-----{model_name}-----")
    model.fit(X_train, y_train)
```

```

y_pred = model.predict(X_test)
y_pred_proba = model.predict_proba(X_test)[: , 1] if hasattr(
    model, "predict_proba") else None

model_results = {
    'Model': model_name,
    'Accuracy': accuracy_score(y_test, y_pred),
    'Precision': precision_score(y_test, y_pred, zero_division=0),
    'Recall': recall_score(y_test, y_pred, zero_division=0),
    'F1 Score': f1_score(y_test, y_pred, zero_division=0),
    'ROC AUC': roc_auc_score(y_test, y_pred_proba) if y_pred_proba \
        is not None else 'N/A',
    # Convert to list for easier JSON serialization
    'Confusion Matrix': confusion_matrix(y_test, y_pred).tolist(),
    'Classification Report': classification_report(y_test, y_pred, zero_division=0)
}
results.append(model_results)

# Store ROC curve data
if y_pred_proba is not None:
    fpr, tpr, _ = roc_curve(y_test, y_pred_proba)
    roc_curves[model_name] = (fpr, tpr, auc(fpr, tpr))

```

```

-----Logistic Regression-----
-----Decision Tree-----
-----Random Forest-----
-----Gradient Boosting-----
-----AdaBoost-----
-----XGBoost-----
-----Support Vector Machine-----
-----K-Nearest Neighbors-----
-----Naive Bayes-----
-----Bagging-----
-----Voting-----

```

In [22]: `# Display results as a DataFrame`
`results_df = pd.DataFrame(results)`
`print(results_df)`

	Model	Accuracy	Precision	Recall	F1 Score	ROC AUC	\
0	Logistic Regression	0.777732	0.757802	0.818182	0.786835	0.852694	
1	Decision Tree	0.972136	0.969084	0.975547	0.972305	0.972126	
2	Random Forest	0.965155	0.949271	0.983035	0.965858	0.995458	
3	Gradient Boosting	0.847363	0.822852	0.886393	0.853441	0.935849	
4	AdaBoost	0.798205	0.772548	0.846847	0.807993	0.888262	
5	XGBoost	0.921511	0.943303	0.897391	0.919775	0.982512	
6	Support Vector Machine	0.817622	0.778986	0.888265	0.830044	0.893296	
7	K-Nearest Neighbors	0.923388	0.874135	0.989704	0.928336	0.964757	
8	Naive Bayes	0.754854	0.737237	0.794080	0.764603	0.816560	
9	Bagging	0.983809	0.987504	0.980110	0.983793	0.995740	
10	Voting	0.875345	0.843937	0.921844	0.881172	0.963457	

```

Confusion Matrix \
0 [[6265, 2235], [1554, 6993]]
1 [[8234, 266], [209, 8338]]
2 [[8051, 449], [145, 8402]]
3 [[6869, 1631], [971, 7576]]
4 [[6369, 2131], [1309, 7238]]
5 [[8039, 461], [877, 7670]]
6 [[6346, 2154], [955, 7592]]
7 [[7282, 1218], [88, 8459]]
8 [[6081, 2419], [1760, 6787]]
9 [[8394, 106], [170, 8377]]
10 [[7043, 1457], [668, 7879]]

```

0	precision	recall	f1-score	...
1	precision	recall	f1-score	...
2	precision	recall	f1-score	...
3	precision	recall	f1-score	...
4	precision	recall	f1-score	...
5	precision	recall	f1-score	...
6	precision	recall	f1-score	...
7	precision	recall	f1-score	...
8	precision	recall	f1-score	...
9	precision	recall	f1-score	...
10	precision	recall	f1-score	...

```
In [23]: # Identify the best model based on a selected metric, e.g., F1 Score
best_model = results_df.loc[results_df['F1 Score'].idxmax()]
```


```
In [24]: print("Best Model Based on F1 Score:")
print(best_model)
```

```
Best Model Based on F1 Score:
Model                                Bagging
Accuracy                             0.983809
Precision                             0.987504
Recall                               0.98011
F1 Score                             0.983793
ROC AUC                              0.99574
Confusion Matrix                     [[8394, 106], [170, 8377]]
Classification Report                precision  recall  f1-score  ...
Name: 9, dtype: object
```

```
In [25]: # Print the metrics of the best model
print(f"\nModel: {best_model['Model']}")
print(f"Accuracy: {best_model['Accuracy']}")
print(f"Precision: {best_model['Precision']}")
print(f"Recall: {best_model['Recall']}")
print(f"F1 Score: {best_model['F1 Score']}")
print(f"ROC AUC: {best_model['ROC AUC']}")
```

```
Model: Bagging
Accuracy: 0.9838094679415733
Precision: 0.9875044206059177
Recall: 0.9801099801099801
F1 Score: 0.9837933059307105
ROC AUC: 0.9957402597402597
```

```
In [26]: # Plot confusion matrix for the best model
best_model_cm = best_model['Confusion Matrix']
plt.figure(figsize=(6, 4))
sns.heatmap(best_model_cm, annot=True, fmt='d', cmap='Blues')
plt.title(f'Confusion Matrix: {best_model["Model"]}')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.show()
```



```
In [ ]: # Print classification report for the best model
print(f"Classification Report for {best_model['Model']}: \n
      {best_model['Classification Report']}")
```

```
Classification Report for Bagging:
              precision    recall  f1-score   support

     0               0.98       0.99       0.98        8500
     1               0.99       0.98       0.98        8547


 accuracy               0.98
macro avg               0.98       0.98       0.98        17047
weighted avg           0.98       0.98       0.98        17047
```

```
In [30]: # Plot confusion matrix for each model
for model_name in models.keys():
    cm = next(item['Confusion Matrix'] for item in results if
              item['Model'] == model_name)

    plt.figure(figsize=(6, 4))
    sns.heatmap(cm, annot=True, fmt='d', cmap='Blues')
    plt.title(f'Confusion Matrix: {model_name}')
    plt.xlabel('Predicted')
    plt.ylabel('Actual')
    plt.show()

    # Print classification report
    report = next(item['Classification Report'] for item in results
                  if item['Model'] == model_name)
    print(f'Classification Report for {model_name}: \n{report}')
```


Confusion Matrix: Logistic Regression



Classification Report for Logistic Regression:

	precision	recall	f1-score	support
0	0.80	0.74	0.77	8500
1	0.76	0.82	0.79	8547
accuracy			0.78	17047
macro avg	0.78	0.78	0.78	17047
weighted avg	0.78	0.78	0.78	17047


Confusion Matrix: Decision Tree



Classification Report for Decision Tree:

	precision	recall	f1-score	support
0	0.98	0.97	0.97	8500
1	0.97	0.98	0.97	8547


accuracy			0.97	17047
macro avg	0.97	0.97	0.97	17047
weighted avg	0.97	0.97	0.97	17047



Classification Report for Random Forest:

	precision	recall	f1-score	support
0	0.98	0.95	0.96	8500
1	0.95	0.98	0.97	8547


accuracy			0.97	17047
macro avg	0.97	0.97	0.97	17047
weighted avg	0.97	0.97	0.97	17047



Classification Report for Gradient Boosting:


	precision	recall	f1-score	support
--	-----------	--------	----------	---------

	0	0.88	0.81	0.84	8500
	1	0.82	0.89	0.85	8547
accuracy				0.85	17047
macro avg		0.85	0.85	0.85	17047
weighted avg		0.85	0.85	0.85	17047




Classification Report for AdaBoost:

	precision	recall	f1-score	support
0	0.83	0.75	0.79	8500
1	0.77	0.85	0.81	8547
accuracy			0.80	17047
macro avg	0.80	0.80	0.80	17047
weighted avg	0.80	0.80	0.80	17047



Classification Report for XGBoost:


	precision	recall	f1-score	support
0	0.90	0.95	0.92	8500
1	0.94	0.90	0.92	8547
accuracy			0.92	17047
macro avg	0.92	0.92	0.92	17047
weighted avg	0.92	0.92	0.92	17047



Classification Report for Support Vector Machine:

	precision	recall	f1-score	support
0	0.87	0.75	0.80	8500
1	0.78	0.89	0.83	8547


accuracy			0.82	17047
macro avg	0.82	0.82	0.82	17047
weighted avg	0.82	0.82	0.82	17047



Classification Report for K-Nearest Neighbors:

	precision	recall	f1-score	support
0	0.99	0.86	0.92	8500
1	0.87	0.99	0.93	8547


accuracy			0.92	17047
macro avg	0.93	0.92	0.92	17047
weighted avg	0.93	0.92	0.92	17047



Classification Report for Naive Bayes:


	precision	recall	f1-score	support
--	-----------	--------	----------	---------

	0	0.78	0.72	0.74	8500
	1	0.74	0.79	0.76	8547
accuracy				0.75	17047
macro avg		0.76	0.75	0.75	17047
weighted avg		0.76	0.75	0.75	17047



Classification Report for Bagging:

	precision	recall	f1-score	support
0	0.98	0.99	0.98	8500
1	0.99	0.98	0.98	8547
accuracy			0.98	17047
macro avg	0.98	0.98	0.98	17047
weighted avg	0.98	0.98	0.98	17047



Classification Report for Voting:


	precision	recall	f1-score	support
0	0.91	0.83	0.87	8500
1	0.84	0.92	0.88	8547
accuracy			0.88	17047
macro avg	0.88	0.88	0.88	17047
weighted avg	0.88	0.88	0.88	17047

```
In [31]: # Define a color palette
colors = sns.color_palette('husl', n_colors=len(roc_curves))
```

```
In [32]: # Plot ROC curves for all models
plt.figure(figsize=(12, 8))
for (model_name, (fpr, tpr, roc_auc)), color in zip(roc_curves.items(), colors):
    plt.plot(fpr, tpr, label=f'{model_name} (AUC = {roc_auc:.2f})', color=color)

plt.plot([0, 1], [0, 1], 'k--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curves for All Models')
plt.legend(loc='lower right')
plt.show()
```

ROC Curves for All Models



```
In [33]: # Assign colors to each model based on the order in results_df
model_colors = {model: color for model, color in zip(results_df['Model'], colors)}
```

```
In [34]: # Bar plot of model accuracies with the same color for each corresponding model
plt.figure(figsize=(12, 8))
sns.barplot(x='Accuracy', y='Model', data=results_df.sort_values(by='Accuracy',
                                                                ascending=False), palette=model_colors)

plt.title('Model Accuracy Comparison')
plt.xlabel('Accuracy')
plt.ylabel('Model')
plt.show()
```

C:\Users\muhammad.abdullah.ai\AppData\Local\Temp\ipykernel_2416\793954673.py:3: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `y` variable to `hue` and set `legend=False` for the same effect.

```
sns.barplot(x='Accuracy', y='Model', data=results_df.sort_values(by='Accuracy',
```

Model Accuracy Comparison

